





Plan of Action

Wednesday, 19 August 2026 1:23 PM



Count based vs Judgement based metrics

Wednesday, 19 August 2026 3:53 PM



X

Variance

complex
helpfulness

4-5
4 eval

You are evaluating whether an AI's answer is correct.

You will be given:

- QUESTION: the question that was asked
- EXPECTED ANSWER: the correct reference answer
- ACTUAL ANSWER: the answer the AI actually gave

Compare the ACTUAL ANSWER against the EXPECTED ANSWER and decide

→ TAG

Compare the ACTUAL ANSWER against the EXPECTED ANSWER and decide how factually correct it is.

Give a score from 0 to 10, where:

- 10 = fully correct
- 0 = completely wrong

QUESTION: {question}

EXPECTED ANSWER: {expected_answer}

ACTUAL ANSWER: {actual_answer}

Score:

0-10

7 → 10% 51%
8 → 51% 40%
6 → 9% 9%
8 7
0-10

↓
g.m
and

8, 7, 6

correcting → rag chatbot

→ golden dataset

LLM
judge


15

ques

correct and

CoT

gpt-4 mode



10,000

steps

Output token.

top-k →

System prompt

LLM / transformer

is 0 → 0.61

the 0 → 0.31

o 0 → 0.01

herf 0 → 0.001

0

0 —

1 —

2 —

3 —

:

:

10

with —

Token	Probability
"8"	0.70
"7"	0.20
"9"	0.05
"the"	0.01
"."	0.01



Score	
8	
7	
9	

These sum to $0.70 + 0.20 + 0.05 = 0.95$.

Divide each by 0.95:

- $8 \rightarrow 0.70 / 0.95 = 0.7368$
- $7 \rightarrow 0.20 / 0.95 = 0.2105$
- $9 \rightarrow 0.05 / 0.95 = 0.0526$

log prob

threshold
0.7
fail

- $8 \times 0.7368 = 5.8947$
- $7 \times 0.2105 = 1.4737$
- $9 \times 0.0526 = 0.4737$

Sum = $5.8947 + 1.4737 + 0.4737 = 7.84$

$7.84 / 10 = 0.784$

0-1
0.

0.7

deepeval imple GEval

```
from deepeval import evaluate
from deepeval.test_case import LLMTestCase, LLMTestCaseParams
from deepeval.metrics import GEval
```

```
correctness = GEval(
```

```
    name="Correctness",
```

```
    criteria="Determine whether the actual output is factually correct based on the expected output.",
```

```
    evaluation_params=[
```

```
        LLMTestCaseParams.INPUT,
```

```
        LLMTestCaseParams.ACTUAL_OUTPUT,
```

```
        LLMTestCaseParams.EXPECTED_OUTPUT,
```

```
    ],
```

```
    model="gpt-4o-mini",
```

```
    threshold=0.7,
```

```
)
```

```
test_case = LLMTestCase(
```

```
    input="What is offline eval?",
```

```
    actual_output="Offline eval runs before deployment.",
```

```
    expected_output="Offline eval is evaluation done before deployment using a fixed golden dataset.",
```

```
)
```

```
evaluate(test_cases=[test_case], metrics=[correctness])
```


built-in metrics

breakdown

prompt

Correctness / faithfulness

Wednesday, 19 August 2020 1:24 PM



✓

x

x

✓

x

x

ans grounded
context

human
expert

golden dataset →
question | correct
answer

—

—

—

—

—

—

—

—

eval pipe line

→ g eval →

Completeness

Wednesday, 19 August 2026 1:24 PM

